

---

## Transaction Management

**Transaction:** Series of actions (reads and writes ) carried out by single user which accesses or changes Database, must be treated as an indivisible unit – Atomic, Logical unit of Work

Transaction transforms Database from one consistent state to another consistent state

Consistent = *Acting or done in the same way over time, especially so as to be fair or accurate*

By default, MySQL runs with *autocommit* mode enabled, i.e. all changes to a table take effect immediately, session variable

Select @@autocommit;   => 1

Transaction:

**Start | Begin Transaction,**

**autocommit;   => 0**

Read (X)

$X = X - N$

Write(X)

Read(Y)

$Y = Y + N$

Write(Y)

**Commit | Rollback   => write out to disk**

**autocommit;   => 1**



## Four properties of a transaction (A.C.I.D. properties)

- **Atomicity** : Transaction is atomic, it is either performed entirely or not performed at all
- **Consistency** : Transaction will transform the Database from one consistent state to another

Transaction won't violate declared system integrity constraints

Database CAN be in a inconsistent state during a transaction, but at the end of the transaction the database is made consistent again.

- **Independence (Isolation)**: Transactions execute independently of one another

Ideally, all the updates from transaction A are not visible to transaction B until transaction A commits

Balance between isolation level and performance.

- **Durability ( Persistence )**: Effects of a successfully completed transaction permanently recorded in the database but not necessarily in database on disk



Data structures that implement each of the ACID properties

- **Atomicity** : *Autocommit, Start / Begin Tx, Commit, Rollback*
- **Consistency** : *Write buffers, Log Files, Crash recovery*
- **Independence (Isolation)**: *Autocommit, Set Isolation Level, Locking*
- **Durability ( Persistence )**: *Write buffers, Replication, UPS, Backups, OS, log files*

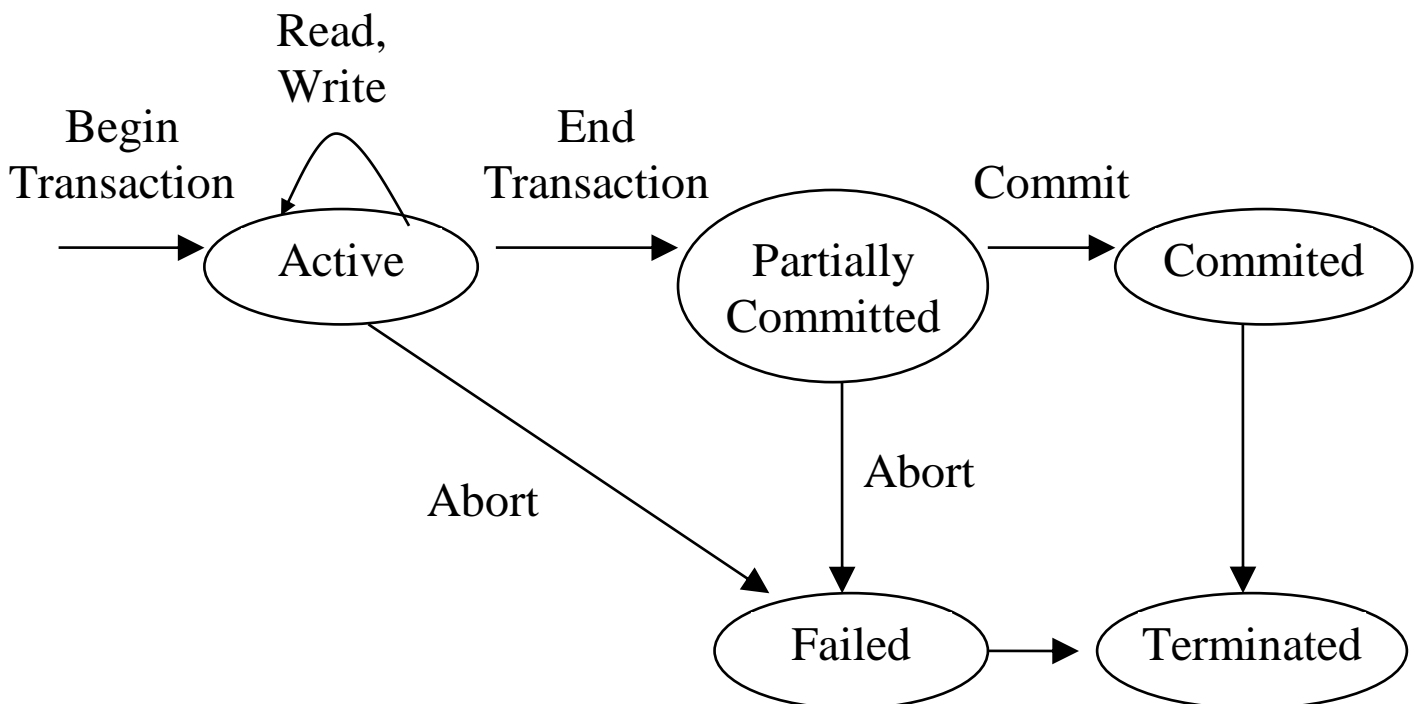


Transaction can have one of two outcomes

1. **Committed**: Database reaches new consistent state
2. **Aborted**: Database restored to previous consistent state- Rollback, Committed transaction cannot be aborted

Transaction Life Cycle, can be in any number of states

- Begin\_Transaction
- Read or Write
- End\_Transaction : Check Database for consistency, see if transaction interfered with other transaction
- Commit\_Transaction : Successful end of transaction
- Failed : Unsuccessful end of transaction, one of checks fails or transaction aborted



## Transaction Schedules

A schedule is the order in which Transactions are executed in a system

### 1. Serial Schedules - no-interleaving

Transactions are executed in a serial manner i.e., no transaction starts until a running transaction has ended.

**Fact:** Nothing Can Go Wrong If The System Executes Transactions Serially

Highest level of isolation but poorest performance!!  
Highest level of consistent data

**On average**, poor performance :

Excellent performance if your Tx at start of queue,  
Very poor performance if your Tx at the end of the queue

Serial schedules are guaranteed to avoid interference and keep the database consistent

### 2. Serializable Schedule - interleaving

Transactions run **concurrently** and are **interleaved**

Transactions possibly executed non-serially but the overall effect on the database is equivalent to a serial schedule



**Serializability:** Set of transactions executing concurrently (i.e. with interleaving), overall effect on database is as though all committed transaction were executed in serial manner, i.e one after another after another

The concurrent execution of a number of transactions must be **Serially Equivalent**



Consider two transaction:

(a) Read (X)

$X = X - N$

Write(X)

Read(Y)

$Y = Y + N$

Write(Y)

(b) Read (X)

$X = X + M$

Write (X)

E.g. Amount (i.e. X) is 100 to start , N=20, M=30

**Serial Schedule: Run all of T1, then run T2**

T1	T2	T1	T2
Read (X)		100	
$X = X - N$		80	
Write(X)		80	
Read(Y)			
$Y = Y + N$			
Write(Y)			
Commit			
	Read (X)		80
	$X = X + M$		110
	Write (X)		110
	Commit		



E.g. Amount (i.e. X) is 100 to start , N=20, M=30

### Serializable Schedule

T1	T2	T1	T2
Read (X)		100	
$X = X - N$		80	
Write(X)		80	
Temporarily Suspend	Read (X)		80
	$X = X + M$		110
	Write (X)		<b>110</b>
Read(Y)			
$Y = Y + N$			
Write(Y)			

T1	T2	T1	T2
	Read (X)		100
	$X = X + M$		130
	Write (X)		130
Read (X)		130	
$X = X - N$		110	
Write(X)		<b>110</b>	
Read(Y)			
$Y = Y + N$			
Write(Y)			

The two Serializable Schedules above resulted in the same consistent effect on the database (i.e. same result as a Serial Schedule) but...this might not always be the case





If several users access a database concurrently, problems may occur

Concurrency problems include:

1. **Lost Updates**: If two Tx write the same data, the second Tx will overwrite the first data value, thereby losing the first Tx's write – Getting seat on flight in labs!!
2. **Dirty Read** - Tx reads a record that has not yet been committed - *reading tentative data*.
3. **Non-Repeatable Read** - Tx re-reads a particular row and finds that the data value has changed or been deleted for that one row

Not always a bad thing!!!

When data has changed, you have to decide whether you want consistent data or concurrent data.

If you need consistent data, then a non-repeatable read is detrimental. If you need current data, then a non-repeatable read is wonderful.

4. **Phantom Read** : In the course of a Tx, two **identical queries** are executed, and the rows returned by the second query are different from the first query.

Not always a bad thing!!!



## 1. Lost Update Problem :

Two transactions that access the **same** Database item/s have their operations interleaved in such a way as to make the value of a Database item incorrect, an update to the database gets lost!!

E.g.

T1	T2		
Read (X)		100	
$X = X - N$		80	
Temporarily Suspend			
	Read (X)		100
	$X = X + M$		130
	Temporarily Suspend		
Write(X)		<b>80</b>	
Read(Y)			
Temporarily Suspend			
	Write (X)		<b>130</b>
$Y = Y + N$			
Write(Y)			

Final value of X is incorrect as T2 reads the value of X before T1 writes it to the Database

Updated value from T1(80) is lost



## 2. Dirty Read (Temporary Update) Problem :

A transaction updates/writes Database item and then the transaction fails for some reason but another transaction has read the updated item

E.g.			
T1	T2	T1	T2
Read (X)		100	
$X = X - N$		80	
Write(X)		80	
Temporarily Suspend			
	<i>Read (X)</i>		80
	<i>Dirty Read</i>		
	$X = X + M$		110
	Write (X)		110
Read(Y)			
<b>T1 Fails</b>			

T1 fails and must change the value of X back to its old value, but T2 has read the temporarily incorrect value of X

Value of item X read by T2 is called Dirty Read as it has been created by an uncommitted transaction

*Reading tentative data* but..... if T1 didn't fail, this would be wonderful as T2 would be reading the most up to date data



### 3. Non-Repeatable Read Problem :

A transaction Reads a particular row in a table, another transaction updates this row

The original transaction re-reads the same row and gets different data

Inconsistent read

E.g.			
T1	T2	T1	T2
Read (X)		<b>100</b>	
$X = X - N$		80	
Write(X)		80	
Temporarily Suspend			
	Read (X)		<b>80</b>
	$X = X + M$		110
	Write (X)		<b>110</b>
Read(X)		<b>110</b>	

Reads a particular row and finds that the data has been changed or deleted for that one row, data was 100, now 110

Not always a bad thing!!!



#### 4. Phantom Read:

T1 executes the same query twice and a different set of rows were returned the second time

E.g.

T1	T2
Select * from users where age between 10 and 30;  return <b>2</b> records. Temporarily Suspend	
Select * from users where age between 10 and 30;  return <b>3</b> records.	Insert into users values ( 3, 'Bob', 27 );  Commit;

Not always a bad thing!!!

Each of the above problems may or may not be an issue for your database

One of the considerations is whether consistent data is your priority or whether concurrent data is your priority



## Independence (Isolation)

Ideally, one Tx cannot interfere with another Tx.

Isolation level

- Defines when changes made by one Tx become visible to other concurrent Tx
- Balance Between Performance and Reliability, Consistency
- Various levels are not "bad" or "good", it depends on what your application requires in terms of isolation.

Various Transaction Isolation levels (**lowest** to **highest**)

1. **Read Uncommitted**
2. **Read Committed**
3. **Repeatable Read** - Default for Innodb
4. **Serrializable** - Default for MySQL

<b>Isolation Level</b>	<b>Dirty Read</b>	<b>Non-repeatable Read</b>	<b>Phantom Read</b>
<b>Read Uncommitted</b>	Allowed	Allowed	Allowed
<b>Read Committed</b>	Prevented	Allowed	Allowed
<b>Repeatable Read</b>	Prevented	Prevented	Allowed
<b>Serializable</b>	Prevented	Prevented	Prevented



1. **Read Uncommitted** –

- Tx B will see the changes that Tx A makes prior to TxA committing
- Great for concurrency / performance provided .... Tx A doesn't do a rollback
- Possibility of **Dirty Reads**, **Phantom reads** and **Non-Repeatable Reads**

2. **Read Committed** –

- Tx B will only see the changes that Tx A makes after TxA commits
- Possibility of **Non-Repeatable Reads** and **Phantom Reads** - Data getting changed in current transaction by other transactions.
- Does not allow dirty reads

3. **Repeatable Read** - Default for Innodb

- **Snapshot** – Version | Copy of database
- Guarantees to see same data for Selects
- Possibility of **Phantom Reads** - Data gets changed in current transaction by another transactions doing Updates/ Insert/Delete – Phantom Read
- <https://bugs.mysql.com/bug.php?id=63870>

4. **Serializable** - Default for SQL

- **Select** gets converted to Read Lock,
- **Insert /Update /Delete** gets converted to Wr lock, Database will lock data until Tx commits
- Lots of locking – Big performance hit!!
- Performs all transactions in sequence, serially



**Begin;**

```
Select * From T;  
Waitfor Delay '00:01:00'  
Select * From T;
```

**Commit;**

- **READ UNCOMMITTED**, second SELECT may return different data. A concurrent transaction may update the record, delete it, insert new records. The second select will **see** the very latest data from all active transactions.
- **READ COMMITTED**, second SELECT may return different data. A concurrent transaction may update the record, delete it, insert new records. The second select will always **see** the very latest data from all committed transactions.
- **REPEATABLE READ**, second SELECT is guaranteed to see the same rows as the first select

New rows may be added by a concurrent transaction during the one minute, but the existing Tx cannot see them.

- **SERIALIZABLE**, second SELECT is guaranteed to see exactly the same rows as the first Select

**Select** automatically acquires read lock in this mode, so no row can be modified/deleted by a concurrent transaction.





## Repeatable Read

Consistent read view or consistent snapshot.

Time	T1	T2
1	<b>Start Tx;</b>	
2	<b>SnapShot</b>	<b>Start Tx;</b>
3	Select * from T;	<b>Snapshot</b>
4	{empty set}	
5		Insert into T values (1, 2);
6	Select * from T;	
7	{empty set}	
8		Select * from T;
9		{1,2}
10		<b>Commit;</b>
11	Select * from T;	
12	{empty set}	
13		
14	<b>Commit;</b>	
15	Select * from T;	
16	{1 ,2}	



## Repeatable Read

Initial value of field x in table is 10

Time	Tx A	Tx B
1	Start Tx;	
2	<b>SnapShot</b>	Start Tx;
	Select x from table; (10)	<b>SnapShot</b>
4	update Table set x=x+10; (20)	
5		select x from Table; (10)
6	Select x from table; (20)	
7	Commit;	
8		select x from Table; (10)
9		Commit;
10		select x from Table; (20)



## Repeatable Read Example of Phantom Read

Time	Tx A	Tx B
1	Start Tx;	
2	<b>SnapShot</b>	Start Tx;
3	Select * from table; (1,1),(2,2),(3,3)	<b>SnapShot</b>
4		select * from Table; (1,1),(2,2),(3,3)
5		Insert into table values (4,4);
6	Select * from table; (1,1),(2,2),(3,3)	
7		Commit;
8	Select * from table; (1,1),(2,2),(3,3)	
9	<b><u>Update</u></b> table Set field2=5; <b>Four rows changed;</b>	
10	Select * from table; (1,5),(2,5),(3,5), (4,5) <b>Phantom Read</b>	



## Read Committed

Set Session Transaction Isolation Level Read Committed;  
No Snapshot here

### Example of Non-Repeatable Read

Time	Tx A	Tx B
1	Start Tx;	
2	select * from Table; (1,1),(2,2),(3,3)	Start Tx;
3	<b>Select * from table where field1=1; (1,1)</b>	
4		select * from Table; (1,1),(2,2),(3,3)
5		<b><u>Update</u></b> table Set field2=5 where field1=1;
6		Commit;
7	<b>Select * from table where field1=1; (1,5) Non-Repeatable Read</b>	



## Concurrency Control

Concurrency Control is managing simultaneous transactions in a database without each transaction interfering with one another

Coordinating the actions of transactions that

- Operate in parallel
- Access **shared** data
- **Potentially** interfere with each other

in order to lead to better transaction throughput and response time

DBMS behaves as if it executes operations sequentially, i.e., one at a time, one after another but .....DBMS typically will execute operations *concurrently* i.e. executing more than one operation at once.

However the DBMS executes transactions, the final and overall effect on the database must be the same as if sequential execution occurred

We will look at three concurrency control techniques :

1. Locking
2. Versioning
3. Multi-version concurrency control (MVCC)



## 1. Locking

Traditional concurrency control

Possible to lock data at different levels:

- **Database** : Entire Database is locked and unavailable  
Used for system backups
- **Table** : Entire Table is locked  
Used for general updates
- **Block or Page** : Physical Page is locked  
Generally not desirable
- **Row Level** : Only requested row is locked  
Most common implementation  
All other rows free,  
But overhead at run-time
- **Field Level** : Only requested field (column) is locked  
Used when an update is to one or two  
fields in a table  
Rarely used as they require considerable  
overhead

Locking level also called **Granularity**



- Fine granularity – Field Row, Page – lots of locks, a performance overhead, more concurrency
- Coarse granularity – Database Table – fewer locks but less overhead, less concurrency

InnoDB supports **Multiple Granularity Locking** which permits coexistence of row locks and table locks

Table Intention Locks are used by a transaction to indicate which type of lock (shared or exclusive) it will require later for a row in that table.

There are two types of intention locks used in InnoDB:

- Intention shared (IS): Transaction intends to set a read lock.
- Intention exclusive (IX): transaction intends to set an write lock.

Intention locks rarely block (only LOCK TABLE.  
....WRITE)

Multiple transactions can acquire IX locks. These locks are at the table-level, not the row-level. An IX lock means that the thread holding it intends to update some rows *somewhere* in the table.

IX locks are only intended to block full-table operations



The idea of intention locks is to allow the lock requests to be resolved in order, but not blocking locks and operations that are compatible

Whether a lock is granted or not can be summarised as follows:

- An X lock is not granted if any other lock (X, S, IX, IS) is held.
- An S lock **is not** granted if an X or IX lock is held.  
An S lock **is** granted if an S or IS lock is held.
- An IX lock **is not** granted if an X or S lock is held.  
An IX lock **is** granted if an IX or IS lock is held.
- An IS lock **is not** granted if an X lock is held.  
An IS lock **is** granted if an S, IX or IS lock is held.

	<i>X</i>	<i>IX</i>	<i>S</i>	<i>IS</i>
<i>X</i>	Conflict	Conflict	Conflict	Conflict
<i>S</i>	Conflict	Conflict	Compatible	Compatible
<i>IX</i>	Conflict	Compatible	Conflict	Compatible
<i>IS</i>	Conflict	Compatible	Compatible	Compatible





Two types of locks :

- **Shared locks (S Locks, Read Locks)**

Allow other transaction to read but not update record

Placing a shared lock on a record prevents another user from placing an exclusive lock on that record

Possible for more than one transaction to hold read locks on same data item

- **Exclusive Lock (Write Lock, X Lock)**

Prevents another transaction from reading and writing a record until it is unlocked

Placing an exclusive lock on a record prevents another Tx from placing any lock on that record

Row Locks are automatically obtained by DBMS, without user involvement


Possible, under certain circumstances, to allow 'lock conversion/upgrading ' i.e. upgrade read lock to write lock or downgrade write lock to read lock



## Two-Phase Locking (2PL)

Two phases:

1. **Growing Phase** where transaction acquire locks
  2. **Shrinking Phase** where transaction release locks
- If T wants to read an object, first obtains an S lock.
  - If T wants to modify an object, first obtains X lock.
  - T can release lock during T
  - But if T releases any lock, it cannot acquire new locks!



Possibility of **lost updates** as once a lock released, data item it can be changed by another transaction before the first transaction commits or Rollbacks.

If lock upgrading allowed, upgrading must be performed during expansion phase and downgrading performed during shrinking phase


Lead to development of Strict 2PL



## Strick Two-Phase Locking (2PL)

### Well Formed Transaction

- If T wants to read an object, first obtains an S lock.
- If T wants to modify an object, first obtains X lock.
- Lock acquisition performed at start of T – should prevent deadlock
- Hold all locks until end of transaction.
- Release all locks together at end of transaction



Drawbacks of strict 2PL is **restricts concurrency** as it locks the item until after until TX commits

No problem then with **Lost Updates**



The locking phase usually implemented as a loop where all the locks are acquired one by one.

If any lock is not acquired on the first attempt, the algorithm releases all the locks it had acquired, and starts to try to get all the locks again - "back-off and re-try"

Can lead to **Transaction Starvation**, if a single Transaction never acquires all the locks needed for it to continue execution.

- It occurs if the waiting scheme for locked items is unfair, giving priority to some transactions over others.
- Starvation happens if same transaction is always chosen as victim.
- Tx never progress.

Starvation Avoidance:

- Switch priorities so that every thread has a chance to have high priority.
- Use FIFO order among competing request.



## Deadlock

Two or more Tx are waiting for each other to finish

Tx's are unable to proceed because each is waiting for the other to do something.

Assume two customers accessing a joint account

T1	T2
Request <b>S</b> Lock on Balance	
Read Balance	Request <b>S</b> Lock on Balance
	Read Balance
Request <b>X</b> lock on Balance (denied)	Request <b>X</b> lock on Balance (denied)
( wait )	( wait )

Above sequence results in a situation called **Deadlock**

T1 is waiting for T2 to release the S lock while T2 is waiting for T1 to release the S lock

Impasse that results when two or more transactions have locked a resource and each must wait for the other to unlock that resource



Overcome this problem using:

- **Deadlock Prevention** - User must lock ALL records at the beginning of a transaction

But, difficult to predict what records will be required

Ties up resources for a long time

- **Deadlock Resolution** - Allow deadlock to occur but build in a mechanism for detecting and breaking deadlock.

Generally results in rollback of some transaction

**InnoDB** has automatic deadlock detection algorithm

Will rollback last transaction to receive lock involved in deadlock

If lock placed by a different storage engine, e.g. MyISAM, MySQL will detect long transactions and rollback long transactions



---

<b>Starvation</b>	<b>Deadlock</b>
Starvation happens if same transaction is always chosen as victim.	Two or more Transactions are waiting for each other to finish
It occurs if the waiting scheme for locked items is unfair, giving priority to some transactions over others.	Transactions are unable to proceed because each is waiting for the other to do something.
Avoidance: <ul style="list-style-type: none"><li>•Switch priorities so that every thread has a chance to have high priority.</li><li>•Use FIFO order among competing request.</li></ul>	Avoidance: <ul style="list-style-type: none"><li>•Acquire locks in a predefined order.</li><li>•Acquire locks immediately before starting.</li></ul>
Tx never progress.	Tx's are waiting for each other.



## 2. Versioning

Locking referred to as **pessimistic** concurrency control as each time record required, DBMS locks record so that other programs cannot use it

Versioning takes **Optimistic** approach in assuming that users will not request same data, at same instant in time. If they do, users may only need to read data, not write the data

In strict Versioning, No Locking of data required

Database or table supports different views for the data depending on when access begins.

Other common terms for this are “time travel,” “copy on write,” or “copy on demand.”

Each transaction is restricted to a view of the Database as of the time that transaction started

At commit time, check performed to determine whether conflict occurred, oldest transaction is given priority

If conflict, all newer transactions notified of the conflict, rolled back and restarted

But, rollbacks *should* be rare

Increased performance improvement due to greater concurrency





Assume two customers accessing a joint account

T1	T2
Read Balance (1,000)	Read Balance (1,000)
Withdraw 200	Withdraw 500
Commit	Commit Return to database Conflict Identified  ... Rollback  Restart Transaction

At commit time, Database discovers that T2's update conflicts, transaction is aborted and restarted using new updated balance



### 3. Multi-Version Concurrency Control (MVCC)

Until recently, nearly all databases used **pessimistic** concurrency model that locked all or portions of the database while a single transaction updated it

MVCC is an **optimistic** concurrency model

Goal being to combine the best properties of a versioning databases with traditional two-phase locking.

Used my Oracle, PostGre, MySQL, SQLServer, IBM DB2, etc

Combination of versioning and Locking

Locking is implemented mainly at the row level



## MVCC in MySQL

InnoDB keeps information about old versions (values) of rows in a *undo log*

InnoDB adds three fields to each row of each table stored in the database.

1. **TR\_ID** : Id of Tx that modified the row.
2. **PTR** : Points to previous version of that row inside the undo log
3. **ROW\_ID** : Counter that increases as new rows are inserted.

InnoDB storage engine uses primarily row-level locking for improved concurrency – at the price of greater complexity

Good concurrency as a given table can be read and modified by different clients at the same time.

InnoDB does not lock the tables in unless explicitly requested via LOCK TABLES command.

InnoDB normally acquires only intentional shared (IS) or intentional exclusive (IX) modes.



Row-level concurrency properties are as follows:

- Different clients can read the same rows simultaneously.
- Different clients can modify **different** rows simultaneously.
- Different clients cannot modify the same row at the same time. If one Tx modifies a row, otherTx cannot modify the same row until the first Tx completes, i.e. releases the exclusive lock

When another transactions sees this update all depends on the Transaction Isolation evel



## MySQL locks

### MyISAM

- Table level locking e.g.

LOCK TABLES emp **READ | WRITE**

Update emp set salary = salary\*1.2

where deptno =30;

UNLOCK TABLES;

### InnoDB :

InnoDB explicitly handles lock, user has nothing to do

- Row level locking , e.g.

**SELECT** ..... from emp **FOR UPDATE**

.....

**Commit | Rollback**

Exclusive lock on the rows selected

**SELECT** ..... from emp **LOCK IN SHARE MODE**

.....

**Commit | Rollback**

Lots of Read locks on the rows selected

Remember plain Select sets no locks

